

Assignment 4. Peeking Blackjack


Hwanjo Yu
CSED342 - Artificial Intelligence

Contact: TA Byoungwoo Kang (bykang@postech.ac.kr)

Deadline : April 15th 2024 at 2:00 pm.

General Instructions

This assignment has only a programming part.

 This icon means you should write code. You can add other helper functions outside the answer block if you want. Do not make changes to files other than `submission.py`.

You should modify the code in `submission.py` between

```
# BEGIN_YOUR_ANSWER
```

and

```
# END_YOUR_ANSWER
```

Your code will be evaluated on two types of test cases, **basic** and **hidden**, which you can see in `grader.py`. Basic tests, which are fully provided to you, do not stress your code with large inputs or tricky corner cases. Hidden tests are more complex and do stress your code. The inputs of hidden tests are provided in `grader.py`, but the correct outputs are not. To run all the tests, type

```
python grader.py
```

This will tell you only whether you passed the basic tests. On the hidden tests, the script will alert you if your code takes too long or crashes, but does not say whether you got the correct output. You can also run a single test (e.g., `3a-0-basic`) by typing

```
python grader.py 3a-0-basic
```

We strongly encourage you to read and understand the test cases, create your own test cases, and not just blindly run `grader.py`.

The search algorithms explored in the previous assignment work great when you know exactly the results of your actions. Unfortunately, the real world is not so predictable. One of the key aspects of an effective AI is the ability to reason in the face of uncertainty.

Markov decision processes (MDPs) can be used to formalize uncertain situations. In this homework, you will implement algorithms to find the optimal policy in these situations. You will then formalize a modified version of Blackjack as an MDP, and apply your algorithm to find the optimal policy.

Problem 1. Value Iteration (Volcano Crossing)

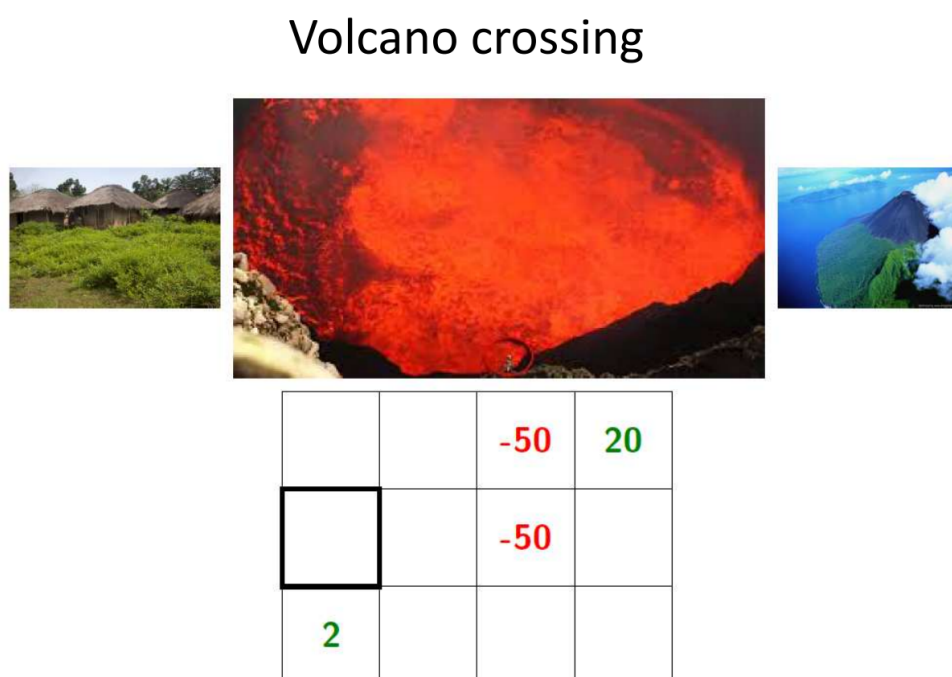


Figure 1: The example image of Volcano Crossing introduced during the class.

In this problem, you will handle an algorithm that performs value iterations for the Volcano Crossing case. There is a grid-world modeled as an M by N grid of states (where M and N are positive integers). The grid world contains at least one volcano grid and one island grid. If the agent moves to any volcano cell, it receives a negative integer reward. If the agent moves to an island cell, it receives a reward of a positive integer. Whenever the agent reaches any island or volcano cell, the travel ends.

To simplify the discussion, we assume that the transition probability $T_{ss'}^a$ is 1 (There is no slip, so the agent will go exactly to the cell it intends to). We only consider the discount factor (which is greater than 0 and less than or equal to 1) and the reward per move (moveReward, which is less than or equal to 0). The agent can move in four directions: east, south, west, and north. If the agent tries to move off the grid, it remains in the same cell.

Problem 1a [5 points]

`VolcanoCrossing` initializes with environment information regarding the grid world (provided as a `numpy.ndarray`), the discount rate, and the `moveReward`. When executing `VolcanoCrossing.value_iteration(n)`, the algorithm performs `n` iterations for the given environment and return the value table containing the value $V_n(s)$ for each state s . To ensure that the value iteration operates correctly, implement `VolcanoCrossing.value_update` to appropriately update the value table.

Problem 2. Blackjack



You will be creating a MDP to describe a modified version of Blackjack. (Before reading the description of the task, first check how `util.ValueIteration.solve` finds the optimal policy of a given MDP such as `util.NumberLineMDP`.)

For our version of Blackjack, the deck can contain an arbitrary collection of cards with different values, each with a given multiplicity. For example, a standard deck would have card values $[1, 2, \dots, 13]$ and multiplicity 4. You could also have a deck with card values $[1, 5, 20]$. The deck is shuffled (each permutation of the cards is equally likely).

The game occurs in a sequence of rounds. Each round, the player either (i) takes the next card from the top of the deck (costing nothing), (ii) peeks at the top card (costing `peekCost`, in which case the card will be drawn in the next round), or (iii) quits the game. (Note: it is not possible to peek twice in a row; if the player peeks twice in a row, then `succAndProbReward()` should return `[]`.)

The game continues until one of the following conditions becomes true:

- The player quits, in which case her reward is the sum of the cards in her hand.
- The player takes a card, and this leaves her with a sum that is strictly greater than the threshold, in which case her reward is 0.
- The deck runs out of cards, in which case it is as if she quits, and she gets a reward which is the sum of the cards in her hand.

In this problem, your state s will be represented as a triple:

`(totalCardValueInHand, nextCardIndexIfPeeked, deckCardCounts)`

As an example, assume the deck has card values $[1, 2, 3]$ with multiplicity 1, and the threshold is 4. Initially, the player has no cards, so her total is 0; this corresponds to state `(0, None, (1, 1, 1))`. At this point, she can take, peek, or quit.

- If she takes, the three possible successor states (each has $1/3$ probability) are

`(1, None, (0, 1, 1))`
`(2, None, (1, 0, 1))`
`(3, None, (1, 1, 0))`

She will receive reward 0 for reaching any of these states.

- If she instead peeks, the three possible successor states are

`(0, 0, (1, 1, 1))`
`(0, 1, (1, 1, 1))`
`(0, 2, (1, 1, 1))`

She will receive reward `-peekCost` to reach these states. From `(0, 0, (1, 1, 1))`, taking yields `(1, None, (0, 1, 1))` deterministically.

- If she quits, then the resulting state will be `(0, None, None)` (note setting the deck to `None` signifies the end of the game).

As another example, let's say her current state is `(3, None, (1, 1, 0))`.

- If she quits, the successor state will be `(3, None, None)`.
- If she takes, the successor states are `(3 + 1, None, (0, 1, 0))` or `(3 + 2, None, None)`. Note that in the second successor state, the deck is set to `None` to signify the game ended with a bust. You should also set the deck to `None` if the deck runs out of cards.

Problem 2a [5 points]

Your task is to implement the game of Blackjack as a MDP by filling out the `succAndProbReward()` function of class `BlackjackMDP`.

Problem 3: Learning to Play Blackjack

So far, we've seen how MDP algorithms can take an MDP which describes the full dynamics of the game and return an optimal policy. But suppose you go into a casino, and no one tells you the rewards or the transitions. We will see how reinforcement learning can allow you to play the game and learn the rules at the same time!

Problem 3a [5 points]

You will first implement a generic Q-learning algorithm `Qlearning`, which is an instance of an `RLAlgorithm`. As discussed in class, reinforcement learning algorithms are capable of executing a policy while simultaneously improving their policy. Look in `util.simulate` to see how the `util.RLAlgorithm` will be used. In short, your `Qlearning` will be run in a simulation of the MDP, and will alternately be asked for an action to perform in a given state (`Qlearning.getAction`), and then be informed of the result of that action (`Qlearning.incorporateFeedback`), so that it may learn better actions to perform in the future.

We are using Q-learning with function approximation, which means $\hat{Q}_{\text{opt}}(s, a) = \mathbf{w} \cdot \phi(s, a)$, where in code, `w` is `self.weights`, ϕ is the `featureExtractor` function, and \hat{Q}_{opt} is `self.getQ`.

We have implemented `Qlearning.getAction` as a simple ϵ -greedy policy. Your job is to implement `Qlearning.incorporateFeedback`, which should take an (s, a, r, s') tuple and update `self.weights` according to the standard Q-learning update.

Problem 3b [5 points]

Now, you'll implement SARSA, which can be considered as a variation of Q-learning. Your task is to fill in `incorporateFeedback` of `SARSA`, which is same with `Qlearning` except the update equation.

Problem 3c [5 points]

Now, we'll incorporate domain knowledge to improve the generalization of RL algorithms for `BlackjackMDP`. Your task is to implement `blackjackFeatureExtractor` as described in the code comments in `submission.py`. This way, the RL algorithm can use what it learned about some states to improve its prediction performance on other states. Using the feature extractor, you should be able to get pretty close to the optimum on some large instances of `BlackjackMDP`.